

海豚调度器: <http://42.194.186.185:12345/dolphinscheduler>

HBase: <http://42.194.186.185:16010/master-status>

脚本

```
1  #/bin/bash
2  #nohup不挂起, &后台运行 , 2>1& > path:不管是正常还是报错的信息都写入
3  #zookeeper启动服务
4  echo '正在启动zookeeper'
5  hosts=(node1 node2 node3)
6  for host in ${hosts[*]}
7  do
8      ssh $host "source /etc/profile;/export/server/zookeeper/bin/zkServer.sh start"
9  done
10
11 #HBase启动服务-- 只需要启动一个
12 echo '正在启动HBase'
13 /export/server/hbase/bin/start-hbase.sh
14
15 #hive的metastore和hiveserver2的服务
16 echo '正在启动hive-metastore'
17 nohup /export/server/hive/bin/hive --service metastore 2>&1 > /tmp/hive-metastore.log &
18 echo '正在启动hive-hiveserver2'
19 nohup /export/server/hive/bin/hive --service hiveserver2 2>&1 > /tmp/hive-
hiveserver2.log &
20 #nohup /export/server/hive/bin/hive --service metastore &
21 #nohup /export/server/hive/bin/hive --service hiveserver2 &
22 #启动MR和spark的历史服务
23 echo '正在启动MR-history'
24 /export/server/hadoop/sbin/mr-jobhistory-daemon.sh start historyserver
25 echo '正在启动spark-history'
26 /export/server/spark/sbin/start-history-server.sh
27 #启动spark连接客户端相关参数
28 echo '正在启动spark客户端连接'
29 /export/server/spark/sbin/start-thriftserver.sh \
30 --hiveconf hive.server2.thrift.port=10001 \
31 --hiveconf hive.server2.thrift.bind.host=node1 \
32 --master local[*]
33
34 #明细定义
35 start-thriftserver.sh \
```

```

36 --name sparksql-thrift-server \
37 --master yarn \
38 --deploy-mode client \
39 --driver-cores 4 \
40
41 --driver-memory 12g \
42 --hiveconf hive.server2.thrift.http.port=10001 \
43 --num-executors 8 \
44 --executor-memory 4g \
45 --conf spark.sql.shuffle.partitions=2
46 #jps -m
47 jps -m
48 #第二种方案(企业)
49 #/export/server/spark/sbin/start-thriftserver.sh \
50 #--hiveconf hive.server2.thrift.port=10001 \
51 #--hiveconf hive.server2.thrift.bind.host=node1 \
52 #--master yarn \
53 #--num-executors 50 \
54 #--executor-cores 4 \
55 #--executor-memory 12G \
56 echo '正在启动海豚调度器'
57 /export/server/dolphinscheduler/bin/start-all.sh
58 echo '查看zookeeper运行状态'
59 /export/server/zookeeper/bin/zkServer.sh status
60
61 echo "正在启动kafka"
62 nohup /export/server/kafka/bin/kafka-server-start.sh
63 /export/server/kafka/config/server.properties 2>&1 > /tmp/kafka.log &
64
65 echo "启动Phoenix客户端"
66 /export/server/phoenixbin/python2 sqlline.py node1:2181
67

```

```

1 kafka
2 #kafka创建主题
3 kafka-topics.sh --create --bootstrap-server node1:9092,node2:9092,node3:9092 --
4 partitions 3 --replication-factor 1 --topic ODS_BASE_LOG_1018
5 #kafka生产者
6 kafka-console-producer.sh --broker-list node1:9092,node2:9092,node3:9092 --topic test001

```

```

6 #kafka消费者
7 kafka-console-consumer.sh --bootstrap-server node1:9092,node2:9092,node3:9092 --topic
  ODS_BASE_LOG_1018
8 #kafka查看主题
9 kafka-topics.sh --bootstrap-server node1:9092,node2:9092,node3:9092 --list
10 #kafka查看主题详情
11 kafka-topics.sh --bootstrap-server node1:9092 --describe --topic test001
12
13 flume
14 #启动flume
15 flume-ng agent -c conf -f /export/server/flume/conf/momo_taildir_source_kafka_sink.conf
  -n a1 -Dflume.root.logger=INFO,console
16
17 phoneix
18 #启动phoneix客户端
19 /export/server/phoenix/bin
  /python2 sqlline.py node1:2181
20
21 启动陌陌jar包
22 java -jar /export/data/momo_init/MoMo_DataGen.jar MoMo_Data.xlsx /export/data/momo_data/
  1000
23

```

## 下载

```

1 #dolphinscheduler调度器
2 wget https://archive.apache.org/dist/dolphinscheduler/1.3.5/apache-dolphinscheduler-
  incubating-1.3.5-dolphinscheduler-bin.tar.gz
3 #MySQL
4 wget https://dev.mysql.com/get/Downloads/MySQL-8.0/mysql-8.0.27-1.el8.x86_64.rpm-
  bundle.tar
5 #spark
6 wget https://archive.apache.org/dist/spark/spark-3.1.2/spark-3.1.2-bin-hadoop3.2.tgz
7 #anaconda3
8 wget https://repo.anaconda.com/archive/Anaconda3-2021.05-Linux-x86_64.sh
9 #hadoop未编译
10 wget https://archive.apache.org/dist/hadoop/common/hadoop-3.3.0/hadoop-3.3.0-src.tar.gz
11 #zookeeper
12 wget https://archive.apache.org/dist/zookeeper/zookeeper-3.7.0/apache-zookeeper-3.7.0-
  bin.tar.gz
13 #hive
14 wget https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz

```

```
15 #sqoop
16 wget https://archive.apache.org/dist/sqoop/1.4.7/sqoop-1.4.7.bin__hadoop-2.6.0.tar.gz
17 #sqoop-commns-lang.jar包
18 wget https://repo1.maven.org/maven2/commons-lang/commons-lang/2.6/commons-lang-2.6.jar
19 #sqoop-hive.jar包
20 wget https://repo1.maven.org/maven2/org/apache/hive/hive-common/3.1.2/hive-common-3.1.2.jar
21 #MySQL8.0jar包
22 wget https://repo1.maven.org/maven2/mysql/mysql-connector-java/8.0.27/mysql-connector-java-8.0.27.jar
23 #MySQL5.0jar包
24 wget https://repo1.maven.org/maven2/mysql/mysql-connector-java/5.1.49/mysql-connector-java-5.1.49.jar
25 #HBase
26 wget https://archive.apache.org/dist/hbase/2.1.0/hbase-2.1.0-bin.tar.gz
27 #Phoenix
28 wget https://archive.apache.org/dist/phoenix/apache-phoenix-5.0.0-HBase-2.0/bin/apache-phoenix-5.0.0-HBase-2.0-bin.tar.gz
29 #kafka
30 wget https://archive.apache.org/dist/kafka/2.4.1/kafka_2.12-2.4.1.tgz
31 #flume
32 wget https://archive.apache.org/dist/flume/1.9.0/apache-flume-1.9.0-bin.tar.gz
```

## python安装库

```
pip install pyarrow,pyspark
```

## python中创建使用pyspark

```
1 import string,time,os
2 from numpy import double
3 from pyspark import StorageLevel
4 from pyspark.sql import SparkSession, Row,functions
5 from pyspark.sql.functions import udf,pandas_udf
6 from pyspark.sql.types import *
7 import pandas as pd
8
9 # 这里选择本地pyspark环境执行spark代码
10 os.environ['JAVA_HOME']=' /export/server/jdk'
11 os.environ['SPARK_HONE'] = ' /export/server/spark'
12 PYSARK_PYTHON = '/root/anaconda3/bin/python'
13 # 当存在多个版本时,不指定很可能会导致出错
```

```

14 os.environ['PYSPARK_PYTHON'] = PYSPARK_PYTHON
15 os.environ['PYSPARK_DRIVER_PYTHON'] = PYSPARK_PYTHON
16 if __name__ == '__main__':
17     # 1.创建SparkSession上下文对象
18     spark = SparkSession.builder.appName('spark_no_hive').master('local[*'])\
19         .config('hive.metastore.uris','thrift://node1:9083')\
20         .config('hive.metastore.warehouse.dir','/user/hive/warehouse')\
21         .config('spark.sql.shuffle.partitions','4')\
22         .enableHiveSupport()\
23         .getOrCreate()
24     spark.sparkContext.setLogLevel('WARN')
25     #Arrow 是一种内存中的列式数据格式，用于spark中以jvm和python进程之间有效的传输数据
26     #需要安装arrow pip install pyarrow
27     #2.开启pyarrow，能加快计算速度，原理2个：1.基于内存减少了序列化和反序列化开销，2.基于向量
    计算vectorize
28     spark.conf.set('spark.sql.execution.arrow.pyspark.enabled','true')

```

## spark写出

```

1 #3.将DataFrame数据保存为text格式
2 #保存为text格式，只能写出一列，不能写出多列
3 df.select(functions.concat_ws('_', 'user_id', 'movieID', 'rating', 'timestamp')).write.mode(
    'overwrite').format('text').save('file:///out1/')
4 #4.将DataFrame数据保存为csv格式 -- header 不忽略字段
5 df.coalesce(1).write.mode('overwrite').format('csv').option('sep', ',').option('header', True).save('file:///csv')
6 #5.将DataFrame数据保存为json格式
7 df.coalesce(1).write.mode('overwrite').format('json').save('file:///json')
8 #6.将DataFrame数据保存为parquet格式
9 df.coalesce(1).write.mode('overwrite').format('parquet').save('file:///parquet')
10 #parquet 语法2 --spark默认支持的格式就是parquet
11 df.coalesce(1).write.mode('overwrite').save('file:///parquet2')
12 #mysql连接
13 df.write.format('jdbc').mode('overwrite').option('url',
    'jdbc:mysql://node1:3306/insurance') \
14     .option('dbtable', 'policy_actuary') \
15     .option('user', 'root').option('password', '123456').save()
16
17 #读取mysql
18 df=spark.read.format('jdbc') \
19     .option('url', 'jdbc:mysql://42.194.186.185:3306/?
    serverTimezone=UTC&characterEncoding=utf8&useUnicode=true') \

```

```

20     .option('dbtable', 'bigdata.tb_top10_movies') \
21     .option('user', 'root').option('password', '123456') \
22     .load()
23 spark=SparkSession.builder.appName("word").master("local[*]").getOrCreate()
24 #2-加载text文件形成DataFrame
25 df=spark.read.format('text').load('path')
26 #简化写法
27 df1=spark.read.text('path')
28
29 #3-加载csv文件形成DataFrame
30 df2=spark.read.format('csv')\
31     .option('sep',';')\
32     .option('header',True)\
33     .option('encoding','utf-8')\
34     .option('inferSchema',True)\
35     .load(path)
36 #4-加载json文件形成DataFrame
37 df3=spark.read.format('json').load('path')
38
39 #5-加载parquet文件形成DataFrame
40 df4=spark.read.format('parquet')\
41     .load(path)
42 #6-添加schema加载文件形成DataFrame
43 df=spark.read.schema('id int,name string, score int')\
44     .csv('path')

```

## 全局

```

1  #JAVA_HOME
2  JAVA_HOME=/export/server/jdk
3  CLASSPATH=.:$JAVA_HOME/lib
4  PATH=$JAVA_HOME/bin:$PATH
5  export JAVA_HOME CLASSPATH PATH
6  #HADOOP_HOME
7  export HADOOP_HOME=/export/server/hadoop
8  export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin
9  #hive和spark的免账号密码登录
10 alias beelinehive="/export/server/apache-hive/bin/beeline -u jdbc:hive2://node1:10000 -n root -p 123456"
11 alias beelinespark="/export/server/spark/bin/beeline -u jdbc:hive2://node1:10001 -n root -p 123456"

```

```
12 #SPARK_HOME
13 export SPARK_HOME=/export/server/spark
14 export PATH=$PATH:$SPARK_HOME/bin
15 #ANACONDA_HOME
16 export ANACONDA_HOME=/root/anaconda3
17 export PATH=$ANACONDA_HOME/bin:$PATH:/root
18 #HIVE
19 export HIVE_HOME=/export/server/hive
20 export PATH=$PATH:$HIVE_HOME/bin
21 #SQOOP_HOME
22 export SQOOP_HOME=/export/server/sqoop
23 export PATH=$PATH:$SQOOP_HOME/bin
24 #ZOOKEEPER_HOME
25 export ZOOKEEPER_HOME=/export/server/zookeeper/
26 export PATH=$PATH:$ZOOKEEPER_HOME/bin
27 #HBASE_HOME
28 export HBASE_HOME=/export/server/hbas
29 export PATH=$PATH:$HBASE_HOME/bin
30 #KAFKA_HOME
31 export KAFKA_HOME=/export/server/kafka
32 export PATH=:$PATH:${KAFKA_HOME}/bin
33 #phoenix_HOME
34 export PHONEIX_HOME=/export/server/phoenix
35 #FLUME
36 export FLUME_HOME=/export/server/flume
37 export PATH=$PATH:$FLUME_HOME/bin
38
```

- 1 【计算链条比较长，每一步骤如果有精度损失，那么叠加到最后，可能会导致结果不准确,则设置如下】
- 2 `set spark.sql.decimalOperations.allowPrecisionLoss=false`
- 3
- 4 【spark和vine导出打印表头，有数据库名】
- 5 `set hive.cli.print.header=true;`
- 6
- 7 【spark和vine导出打印表头，去数据库名】
- 8 `set hive.cli.print.header=true;set hive.resultset.use.unique.column.names=false;`
- 9
- 10 【spark设置shuffle并行度】

```
11 spark.sql('set spark.sql.shuffle.partitions=4')
12 set spark.sql.shuffle.partitions=4
13
14 【hdfs 查看总文件大小】
15 hdfs dfs -du -h path
16
17 【hdfs副本数】
18 dfs.replication
19
20 【修复MySQL元数据，和分区表的实际目录一致】
21 msck repair table 库.表;
22
23 【开启hive中的mapjoin】--大表的数据会分成多个task和整张小表去join
24 set hive.auto.convert.join=true [hive2默认为true]
25
26 【hive和spark的mapjoin小表大小】
27 set hive.auto.convert.join.noconditionaltask.size=1000000; 【默认为10M单位字节】
28 set spark.sql.autoBroadcastJoinThreshold=51200000;
29
30 【hive的bucket mapjoin 分桶】
31 set hive.optimize.bucketmapjoin = true; 【默认为true】
32
33 【sortmerge bucket mapjoin（大表join大表smb）】
34 set hive.auto.convert.join=true;
35 set hive.optimize.bucketmapjoin = true; -- 开启 bucket map join
36 set hive.auto.convert.sortmerge.join=true; -- 开启 SBM join支持--归并
37 set hive.optimize.bucketmapjoin.sortedmerge = true; -- 自动尝试开启 SMB join
38 --set hive.enforce.sorting=true; -- 开启强制排序
39 --set hive.enforce.bucketing=true;-- 注释的只能在配置文件写死或者不执行
40 --set hive.input.format=org.apache.hadoop.hive.ql.io.BucketizedHiveInputFormat;
41 set hive.auto.convert.join.noconditionaltask = true;
42 --下面可以不设置
43 set hive.auto.convert.join.noconditionaltask.size = 50000000;
44 set hive.auto.convert.sortmerge.join.bigtable.selection.policy
45     = org.apache.hadoop.hive.ql.optimizer.TableSizeBasedBigTableSelectorForAutoSMJ;
46
47 【MR的task内存】
48 set mapreduce.map.java.opts=-Xmx6000m;
49 set mapreduce.map.memory.mb=6096;
```



```
50 set mapreduce.reduce.java.opts=-Xmx6000m;
51 set mapreduce.reduce.memory.mb=6096;
52
53 【MR缓冲区大小】
54 set mapreduce.task.io.sort.mb=100
55
56 【MR环形缓冲区溢出的阈值】
57 mapreduce.map.sort.spill.percent=0.8
58
59 【reduce拉取并行度】
60 set mapreduce.reduce.shuffle.parallelcopies=8
61 set mapreduce.reduce.shuffle.read.timeout=180000
62
63 【MR小文件处理】
64 #设置Hive中底层MapReduce读取数据的输入类：将所有文件合并为一个文件作为输入
65 set hive.input.format=org.apache.hadoop.hive.ql.io.CombineHiveInputFormat;
66 #如果hive的程序，只有maptask，将MapTask产生的所有小文件进行合并
67 set hive.merge.mapfiles=true;
68 #如果hive的程序，有Map和ReduceTask,将ReduceTask产生的所有小文件进行合并
69
70 set hive.merge.mapredfiles=true;
71 #每一个合并的文件的大小
72
73 set hive.merge.size.per.task=256000000;
74 #平均每个文件的大小，如果小于这个值就会进行合并
75
76 set hive.merge.smallfiles.avgsize=16000000;
77
78 【查看hive和spark的执行计划】
79 explain select...
80
81 【markdown代码折叠】
82 <details>
83 <summary><b>点击查看完整代码</b></summary>
84 <pre><code>
85 </code></pre>
86 </details>
87
88 【Linux关闭邮件提醒】
89 echo "unset MAILCHECK" >> /etc/profile
90 source /etc/profile
91
```

```
88  【Linux ifconfig没有ip】
89  [方案1推荐]
90  [root@localhost~]# ifup ens33
91  Error: Connection activation failed: No suitable device found for this
    connection(device lo not available because device is strictly unmanaged).
92
93  [root@localhost~]# chkconfig NetworkManager off
94  Note: Forwarding request to 'systemctl disable NetworkManager.service'.
95  Removed symlink /etc/systemd/system/multi-user.target.wants/NetworkManager.service.
96  Removed symlink /etc/systemd/system/dbus-org.freedesktop.NetworkManager.service.
97  Removed symlink /etc/systemd/system/dbus-org.freedesktop.nm-dispatcher.service.
98  Removed symlink /etc/systemd/system/network-online.target.wants/NetworkManager-wait-
    online.service.
99
100 [root@localhost~]# chkconfig network on
101 [root@localhost~]# service NetworkManager stop
102 Redirecting to /bin/systemctl stop NetworkManager.service
103 [root@localhost~]# service network start
104 Starting network (via systemctl):[ ok]
105
106 [方案2]
107 ip add show
108 将link/ether后的MAC地址写到ens33中的HWADDR
109 vim /etc/sysconfig/network-scripts/ifcfg-ens33
110
111  【Linux搜索指定文件类型的指定内容】
112 grep -rin "9820" --include "*.xml" /export/server/
113 find / -name "*.log" | grep "error"
114
115  【SparkSQL报错Filesystem Closed】
116 修改hadoop配置文件core-site.xml:
117 <property>
118     <name>fs.hdfs.impl.disable.cache</name>
119     <value>true</value>
120 </property>
121
122  【启动hive的metastore、hiveserver2服务】
123 nohup hive --service metastore 2>&1 > /tmp/hive-metastore.log &
124 nohup hive --service hiveserver2 2>&1 > /tmp/hive-hiveserver2.log &
125
```

```
126
127 【hive元数据初始化和更新】
128 schematool -dbType mysql -initSchema
129 schematool -dbType mysql -upgradeSchema
130
131 【mysql的授权语句】
132 GRANT ALL PRIVILEGES ON hive.* TO 'root'@'%' IDENTIFIED BY '123456';
133 flush privileges;
134
135 【hive的jdbc连接地址】
136 jdbc:hive2://node1:10000
137
138
139 /export/server/zookeeper/bin/zkServer.sh stop
140 /export/server/zookeeper/bin/zkServer.sh start
141 /export/server/zookeeper/bin/zkServer.sh status
142
143 /export/server/dolphinscheduler/bin/stop-all.sh
144 /export/server/dolphinscheduler/bin/start-all.sh
145
146
147
148 【开启Hive的本地模式】
149 set hive.exec.mode.local.auto=true; (默认为false)
150 【hive开启负载均衡】平均分区，让每个分区都是同样的数据，如果多个去重则会报错，这个针对group by
151 hive.groupby.skewindata=true
152
153
154
155 【hadoop退出安全模式】
156 hdfs dfsadmin -safemode leave
157 hdfs dfsadmin -safemode forceExit
158
159 【重新启动zk和kafka】
160 zkServer.sh stop
161 kafka-server-stop.sh
162 zkServer.sh start
163 nohup /export/server/kafka/bin/kafka-server-start.sh
    /export/server/kafka/config/server.properties &
164
```

```
165 【彻底删除kafka主题】
166 kafka-topics.sh --zookeeper node3:2181 --list
167 kafka-topics.sh --zookeeper node3:2181 --delete --topic spark_kafka
168
169 【zkCli.sh】
170 ls /config/topics
171 rmr /config/topics/spark_kafka
172 rmr /brokers/topics/spark_kafka
173 rmr /admin/delete_topics/spark_kafka
174
175 【创建主题】
176 kafka-topics.sh --zookeeper node3:2181 --create --topic spark_kafka --partitions 3 --
    replication-factor 1
177 kafka-topics.sh --zookeeper node3:2181 --list
178
179 【启动生产者和消费者】
180 kafka-console-producer.sh --broker-list node3:9092 --topic spark_kafka
181 kafka-console-consumer.sh --from-beginning --bootstrap-server node3:9092 --topic
    spark_kafka
182 kafka-console-consumer.sh --from-beginning --bootstrap-server node3:9092 --topic
    __consumer_offsets
183
184
185 【完整的删除再重建主题步骤start
186 关闭生产者和消费者
187
188 kafka-topics.sh --zookeeper node3:2181 --list
189 kafka-topics.sh --zookeeper node3:2181 --delete --topic stationTopic
190
191 kafka-topics.sh --zookeeper node3:2181 --list
192
193 zkCli.sh
194 ls /config/topics
195 rmr /config/topics/stationTopic
196 rmr /brokers/topics/stationTopic
197 rmr /admin/delete_topics/stationTopic
198
199 删除log文件目录
200
201 kafka-topics.sh --zookeeper node3:2181 --create --topic stationTopic --partitions 3 --
    replication-factor 1
```

```
202 kafka-topics.sh --zookeeper node3:2181 --list
203 再次重启kafka
204 kafka-server-stop.sh
205 nohup /export/server/kafka/bin/kafka-server-start.sh
    /export/server/kafka/config/server.properties &
206
207 启动生产者
208 kafka-console-producer.sh --broker-list node3:9092 --topic stationTopic
209 启动消费者
210 kafka-console-consumer.sh --bootstrap-server node3:9092 --topic stationTopic --from-
    beginning
211 完整的删除再重建主题步骤end】
212
213
214 【安装redis】
215 tar
216 yum install gcc
217 cd /redis3.2.8
218 make
219 make PREFIX=/usr/local/src/redis install
220 cp /redis3.2.8/redis.conf /redis/bin/redis.conf
221 启动redis
222 redis-server redis.conf
223 redis-cli -h 192.168.88.163
224
225 【启动spark-thriftserver】
226 /export/server/spark/sbin/start-thriftserver.sh \
227     --hiveconf hive.server2.thrift.port=10001 \
228     --hiveconf hive.server2.thrift.bind.host=node1 \
229     --master local[*]
230
231 /export/server/spark/sbin/start-thriftserver.sh \
232     --hiveconf hive.server2.thrift.port=10001 \
233     --hiveconf hive.server2.thrift.bind.host=p1 \
234     --master local[*]
235
236
237 /export/server/spark/sbin/start-thriftserver.sh \
238     --hiveconf hive.server2.thrift.port=10001 \
239     --hiveconf hive.server2.thrift.bind.host=p1 \
```

```
240 --master yarn \  
241 --deploy-mode client \  
242 --driver-memory 1g \  
243 --executor-memory 2g \  
244 --executor-cores 4 \  
245 --num-executors 25 \  
246  
247 !connect jdbc:hive2://192.168.88.166:10001  
248  
249 【kafka】  
250 启动Zookeeper 服务  
251 zookeeper-daemon.sh start  
252  
253 启动Kafka 服务  
254 kafka-daemon.sh start  
255  
256  
257 【免秘钥登录】  
258 ssh-keygen -t rsa  
259 ssh-copy-id node1  
260 scp /root/.ssh/authorized_keys node2:/root/.ssh  
261 scp /root/.ssh/authorized_keys node3:/root/.ssh  
262  
263 【sqoop】  
264 sqoop import \  
265 -D mapred.job.name=sqoop_import_dd_table \  
266 --connect "jdbc:mysql://192.168.88.163:3306/insurance" \  
267 --username root \  
268 --password "123456" \  
269 --table dd_table \  
270 --hive-import \  
271 --hive-database insurance \  
272 --hive-table dd_table \  
273 --hive-overwrite \  
274 -m 1 \  
275 --fields-terminated-by ',' \  
276 --delete-target-dir  
277  
278 【mysql】
```

```
279 mysqldump insurance --add-drop-trigger --result-file=/opt/insurance/insurance.sql --
    user=root --host=192.168.88.163 --port=3306
280 mysql -uusername -ppassword -h ip_address -P 3306 --default-character-set = utf8mb4 $
    {i} </ data / mysql_back / $ {i} .sql
281 mysql -uroot -p --default-character-set=utf8mb4 insurance </opt/insurance/insurance.sql
282
283
284 【dolphinscheduler】
285 CREATE DATABASE dolphinscheduler DEFAULT CHARACTER SET utf8 DEFAULT COLLATE
    utf8_general_ci;
286 set global validate_password.policy=0;
287 set global validate_password.length=1;
288
289 GRANT ALL PRIVILEGES ON dolphinscheduler.* TO '{user}'@'%' IDENTIFIED BY '{password}';
290 flush privileges;
291
292 安装dolphinscheduler 要启动zookeeper
293 一定要配在ds的env中配置sqoop的环境变量
294 需要root给租户授权最大。
295 谁将脚本上传到Linux，谁就记得授执行权限。
296
297 安装dos2unix命令
298 dos2unix /opt/insurance/sqoop/sqoop_import_mort_10_13.sh
299
300 修改工作流，要先下线
301
302
303 【structured Streaming】
304 --memory sink
305 CREATE TABLE db_spark.tb_word_count (
306     id int NOT NULL AUTO_INCREMENT,
307     word varchar(255) NOT NULL,
308     count int NOT NULL,
309     PRIMARY KEY (id),
310     UNIQUE KEY word (word)
311 ) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci;
312
313 REPLACE INTO tb_word_count (id, word, count) VALUES (NULL, ?, ?);
314
315
316 【spark yarn Pi】
```

```
317 /export/server/spark/bin/spark-submit \  
318 --master yarn \  
319 --class org.apache.spark.examples.SparkPi \  
320 ${SPARK_HOME}/examples/jars/spark-examples_2.11-2.4.5.jar \  
321 10  
322  
323 【WordCount yarn】  
324 /export/server/spark/bin/spark-submit \  
325 --master yarn \  
326 --driver-memory 512m \  
327 --executor-memory 512m \  
328 --executor-cores 1 \  
329 --num-executors 2 \  
330 --queue default \  
331 --class cn.itcast.spark._2SparkWordCount \  
332 /opt/spark-chapter01-1.0-SNAPSHOT.jar  
333  
334  
335  
336 【 Run application locally on 8 cores】  
337 /export/server/spark/bin/spark-submit \  
338 --class org.apache.spark.examples.SparkPi \  
339 --master local[8] \  
340 ${SPARK_HOME}/examples/jars/spark-examples_2.11-2.4.5.jar \  
341 100  
342  
343 # Run on a Spark standalone cluster in client deploy mode  
344 ./bin/spark-submit \  
345 --class org.apache.spark.examples.SparkPi \  
346 --master spark://207.184.161.138:7077 \  
347 --executor-memory 20G \  
348 --total-executor-cores 100 \  
349 ${SPARK_HOME}/examples/jars/spark-examples_2.11-2.4.5.jar \  
350 1000  
351  
352 # Run on a Spark standalone cluster in cluster deploy mode with supervise  
353 ./bin/spark-submit \  
354 --class org.apache.spark.examples.SparkPi \  
355 --master spark://207.184.161.138:7077 \  
356 --deploy-mode cluster \  

```



```
357 --supervise \  
358 --executor-memory 20G \  
359 --total-executor-cores 100 \  
360 /path/to/examples.jar \  
361 1000  
362  
363 # Run on a YARN cluster  
364 export HADOOP_CONF_DIR=XXX  
365 ./bin/spark-submit \  
366 --class org.apache.spark.examples.SparkPi \  
367 --master yarn \  
368 --deploy-mode cluster \ # can be client for client mode  
369 --executor-memory 20G \  
370 --num-executors 50 \  
371 /path/to/examples.jar \  
372 1000  
373  
374 # Run a Python application on a Spark standalone cluster  
375 ./bin/spark-submit \  
376 --master spark://207.184.161.138:7077 \  
377 examples/src/main/python/pi.py \  
378 1000  
379  
380 # Run on a Mesos cluster in cluster deploy mode with supervise  
381 ./bin/spark-submit \  
382 --class org.apache.spark.examples.SparkPi \  
383 --master mesos://207.184.161.138:7077 \  
384 --deploy-mode cluster \  
385 --supervise \  
386 --executor-memory 20G \  
387 --total-executor-cores 100 \  
388 http://path/to/examples.jar \  
389 1000  
390  
391 # Run on a Kubernetes cluster in cluster deploy mode  
392 ./bin/spark-submit \  
393 --class org.apache.spark.examples.SparkPi \  
394 --master k8s://xx.yy.zz.www:443 \  
395 --deploy-mode cluster \  
396 --executor-memory 20G \  

```

```
397 --num-executors 50 \  
398 http://path/to/examples.jar \  
399 1000  
400
```

```
1 hive动态分区  
2 -----因为采用了动态分区插入技术 因此需要设置相关参数-----  
3 --分区  
4 SET hive.exec.dynamic.partition=true;  
5 SET hive.exec.dynamic.partition.mode=nonstrict;  
6 set hive.exec.max.dynamic.partitions.pernode=10000;  
7 set hive.exec.max.dynamic.partitions=100000;  
8 set hive.exec.max.created.files=150000;  
9 --hive压缩  
10 set hive.exec.compress.intermediate=true;  
11 set hive.exec.compress.output=true;  
12 --写入时压缩生效  
13 set hive.exec.orc.compression.strategy=COMPRESSION;
```

```
1 --开启分桶，如果不开启，是不会启动多个reduce分桶的  
2 set hive.enforce.bucketing=true;  
3 --2.x版本可以通过以下参数禁止使用load语句加载数据到表中  
4 set hive.strict.checks.bucketing = true;
```

常见的文件类型：textfile/orc/parquet

```
1 hiveorc索引  
2 stored as orc ('orc.create.index'='true')  
3 查询数据时，开启row group index 过滤查询  
4 hive.optimize.index.filter=true
```

```
1 hive优化  
2 矢量化查询：开启了：Hive每个批次读取1024条，处理1024条  
3 set hive.vectorized.execution.enabled = true;  
4 set hive.vectorized.execution.reduce.enabled = true;  
5  
6 零拷贝：  
7 功能：数据在操作系统中从内存中不用经过多次拷贝直接读取
```

```
8 不开启：必须在内存中经过多次交换才能读取到数据
9 开启了：数据可以直接从内存中读取
10 set hive.exec.orc.zerocopy=true;
11
12 关联优化器
13 功能：Hive在解析SQL语句，转换为MapReduce时候，可以将相关联的部分合并在一起执行
14 自动判断所有执行过程语法数是否有重合的过程，放在一起执行
15 set hive.optimize.correlation=true;
16
17 动态生成分区的线程数
18 说明：在执行动态分区过程中，可以运行多少个线程来生产分区数据，线程数量越多，执行效率越高，前提是有资源
19 hive.load.dynamic.partitions.thread 默认值为 15
```

## hive建表

```
1  hive建表语句
2  create [external | temporary] table [dbname.]tbname(
3      colName1 type1 comment '',
4      colName2 type1 comment '',
5      colName3 type1 comment '',
6      .....
7      colNameN type1 comment ''
8
9  ) comment '表的注释'
10 partitioned by (col type)
11 clustered by (col) [sorted by col desc] into N buckets
12 row format delimited fields terminated by
13 stored as orc [tblproperties ("orc.compress"="[ZLIB|SNAPPY]")]
14 location '指定路径'
15
16 jdbc:mysql://hadoop01:3306/yipin?useUnicode=true&characterEncoding=UTF-8&autoReconnect=true'
```

## mapreduce内存溢出优化

```
1 MapReduce的内存配置
2 map阶段:mapreduce.map.memory.mb : 每一个map默认申请的内存大小，默认值为 0 表示自动获取；
   mapreduce.map.java.opts : 每一个map的jvm的内存大小；
3 注意： mapreduce.map.java.opts一定要小于mapreduce.map.memory.mb;
4
```

```
5  reduce阶段:mapreduce.reduce.memory.mb : 每一个reduce默认申请的内存大小, 默认值为 0 表示自动
   获取; mapreduce.reduce.java.opts : 每一个reduce的jvm的内存大小;
6  注意:      mapreduce.reduce.java.opts 一定要小于mapreduce.reduce.memory.mb;
7
8  注意: MR中所有的内存配置, 都不能大于nodemanager的内存大小
9
10 jvm-java的格式为: -Xmx4096m
11 其他: 4096
```

## 压缩配置

```
1  yarn配置: 在CM中直接配置
2  mapreduce.map.output.compress 是否开启map端压缩配置      默认开启的
3  mapreduce.map.output.compress.codec  map端采用何种压缩方案
4      建议配置为: org.apache.hadoop.io.compress.SnappyCodec
5
6  mapreduce.output.fileoutputformat.compress 是否开启reduce端压缩配置 默认不开启的 对最终结
   果有影响
7  mapreduce.output.fileoutputformat.compress.codec reduce端需要采用何种压缩操作
8      建议配置为: org.apache.hadoop.io.compress.SnappyCodec
9
10 mapreduce.output.fileoutputformat.compress.type 采用压缩的方式
11      建议: block 块压缩
12
13 hive配置: 此配置需要在会话中执行 MR开了压缩, 开启以下才会进行压缩
14 set hive.exec.compress.intermediate=true; 开启中间结果的压缩
15 set hive.exec.compress.output=true; 是否开启最终结果压缩
16
```