

# 1. 下载Flink

```
1 wget https://archive.apache.org/dist/flink/flink-1.14.4/flink-1.14.4-bin-scala_2.11.tgz
```

- 解压、创建软连接

```
1 tar -zxvf flink-1.14.4-bin-scala_2.11.tgz -C /export/server
2
3 cd /export/server
4 ln -s flink-1.14.4-bin-scala_2.11.tgz flink
```

- 配置文件(修改以下参数)

```
1 cd /export/server/flink/conf
2 vim flink-conf.yaml
3
4 # jobManager 的IP地址（配置HA需要把想要的master设置对应的本机名字，比如node2，下载则需要修改为
  node2）
5 jobmanager.rpc.address: node1
6 # JobManager 的端口号
7 jobmanager.rpc.port: 6123
8 # JobManager JVM heap 内存大小
9 jobmanager.memory.process.size: 1600m
10 # TaskManager JVM heap 内存大小
11 taskmanager.memory.process.size: 1728m
12 # 每个 TaskManager 提供的任务 slots 数量大小
13 taskmanager.numberOfTaskSlots: 4
14 #是否进行预分配内存，默认不进行预分配，这样在我们不使用flink集群时候不会占用集群资源
15 #taskmanager.memory.preallocate: false
16 # 程序默认并行计算的个数
17 parallelism.default: 1
18 #JobManager的Web界面的端口（默认：8081）
19 jobmanager.web.port: 8081
```

- 配置worker文件

```
1 vim worker
2
3 node1
4 node2
5 node3
```

- 配置环境变量

```
1 vim /etc/profile
2 FLINK_HOME=/export/server/flink
3 PATH=$PATH:$FLINK_HOME/bin
4 # 立即生效
5 source /etc/profile
6 scp /etc/profile node2:/etc
7 scp /etc/profile node2:/etc
8 #都需要source一下
```

- 分发

```
1 scp -r /export/server/flink node2:/export/server
2 scp -r /export/server/flink node3:/export/server
```

- 开启集群

```
1 start-cluster.sh
```

- 查看当前Flink集群的状态webui

```
1 node1:8081
```

The screenshot shows the Apache Flink Dashboard interface. On the left is a sidebar with navigation links: Overview, Jobs, Task Managers, Job Manager, and Submit New Job. The main content area is divided into several sections:

- Overview:** Shows 'Available Task Slots' as 6 (labeled '当前整个集群总的 slot 数') and 'Running Jobs' as 0 (labeled '正在执行 job 数'). Below these are 'Total Task Slots: 6' and 'Task Managers: 3'.
- Running Job List:** A table with columns: Job Name, Start Time, Duration, End Time, Tasks, and Status. It is currently empty, showing 'No Data'.
- Completed Job List:** A table with columns: Job Name, Start Time, Duration, End Time, Tasks, and Status. It is also empty, showing 'No Data'.

Handwritten red annotations are present on the image:

- A red box around the 'Overview' tab in the sidebar.
- A red box around the 'Running Jobs' section.
- A red box around the 'Running Job List' table.
- A red box around the 'Completed Job List' table.
- Red circles and lines around the 'Start Time', 'Duration', 'End Time', and 'Status' columns in the 'Completed Job List' table, with handwritten notes: '开了' (opened) near Start Time, '跑了多久' (how long it ran) near Duration, '结束了' (ended) near End Time, and '成功?' (successful?) near Status.
- Red text 'yarn job list' is written across the 'Running Job List' table.
- Red text '提交作业web页面' (submit job web page) is written near the 'Submit New Job' link in the sidebar.

- 执行wordcount任务执行, run word-count案例

```
1 flink run -p 2 /export/server/flink/examples/batch/WordCount.jar
```

- 参数解释

```
1 flink run 提交执行任务 类似于 spark-submit
2 -p 1 并行度设置为1
```

```
3  --input  当前输入的参数
4  /export/server/flink/examples/batch/WordCount.jar  jar包位置
```

## HA高可用部署方式

- 修改配置文件(node1)

```
1  cd /export/server/flink/conf
2  vim flink-conf.yaml
3
4  #####
5  # Licensed to the Apache Software Foundation (ASF) under one
6  # or more contributor license agreements. See the NOTICE file
7  # distributed with this work for additional information
8  # regarding copyright ownership. The ASF licenses this file
9  # to you under the Apache License, Version 2.0 (the
10 # "License"); you may not use this file except in compliance
11 # with the License. You may obtain a copy of the License at
12 #
13 # http://www.apache.org/licenses/LICENSE-2.0
14 #
15 # Unless required by applicable law or agreed to in writing, software
16 # distributed under the License is distributed on an "AS IS" BASIS,
17 # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
18 # See the License for the specific language governing permissions and
19 # limitations under the License.
20 #####
21
22
23 #=====
24 # Common
25 #=====
26
27 # The external address of the host on which the JobManager runs and can be
28 # reached by the TaskManagers and any clients which want to connect. This setting
29 # is only used in Standalone mode and may be overwritten on the JobManager side
30 # by specifying the --host <hostname> parameter of the bin/jobmanager.sh executable.
31 # In high availability mode, if you use the bin/start-cluster.sh script and setup
32 # the conf/masters file, this will be taken care of automatically. Yarn
33 # automatically configure the host name based on the hostname of the node where the
```

```
34 # JobManager runs.
35
36 # jobManager 的IP地址（node2需要改为node2，有几个mater就改几个）
37 jobmanager.rpc.address: node1
38
39 # The RPC port where the JobManager is reachable.
40
41 # JobManager 的端口号
42 jobmanager.rpc.port: 6123
43
44
45 # The total process memory size for the JobManager.
46 #
47 # Note this accounts for all memory usage within the JobManager process, including JVM
  metaspace and other overhead.
48
49 # JobManager JVM heap 内存大小
50 jobmanager.memory.process.size: 1600m
51
52
53 # The total process memory size for the TaskManager.
54 #
55 # Note this accounts for all memory usage within the TaskManager process, including JVM
  metaspace and other overhead.
56
57 # TaskManager JVM heap 内存大小
58 taskmanager.memory.process.size: 1728m
59
60 # To exclude JVM metaspace and overhead, please, use total Flink memory size instead of
  'taskmanager.memory.process.size'.
61 # It is not recommended to set both 'taskmanager.memory.process.size' and Flink memory.
62 #
63 # taskmanager.memory.flink.size: 1280m
64
65 # The number of task slots that each TaskManager offers. Each slot runs one parallel
  pipeline.
66
67 # 每个 TaskManager 提供的任务 slots 数量大小
68 taskmanager.numberOfTaskSlots: 4
69
```

```
70 #是否进行预分配内存，默认不进行预分配，这样在我们不使用flink集群时候不会占用集群资源
71 #taskmanager.memory.preallocate: false
72
73 # The parallelism used for programs that did not specify and other parallelism.
74
75 # # 程序默认并行计算的个数
76 parallelism.default: 1
77
78 # The default file system scheme and authority.
79 #
80 # By default file paths without scheme are interpreted relative to the local
81 # root file system 'file:///'. Use this to override the default and interpret
82 # relative paths relative to a different file system,
83 # for example 'hdfs://mynamenode:12345'
84 #
85 # fs.default-scheme
86
87 #=====
88 # High Availability
89 #=====
90
91 # The high-availability mode. Possible options are 'NONE' or 'zookeeper'.
92 #
93 #开启
94 high-availability: zookeeper
95
96 # The path where metadata for master recovery is persisted. While ZooKeeper stores
97 # the small ground truth for checkpoint and leader election, this location stores
98 # the larger objects, like persisted dataflow graphs.
99 #
100 # Must be a durable file system that is accessible from all nodes
101 # (like HDFS, S3, Ceph, nfs, ...)
102 #
103 # 存储JobManager的元数据到HDFS,用来恢复JobManager 所需的所有元数据
104 high-availability.storageDir: hdfs://node1:8020/flink/ha/
105
106 # The list of ZooKeeper quorum peers that coordinate the high-availability
107 # setup. This must be a list of the form:
108 # "host1:clientPort,host2:clientPort,..." (default clientPort: 2181)
```

```
109 #
110 #zookeeper集群地址
111 high-availability.zookeeper.quorum: node1:2181,node2:2181,node3:2181
112
113
114 # ACL options are based on
115 # https://zookeeper.apache.org/doc/r3.1.2/zookeeperProgrammers.html#sc_BuiltinACLSchemes
116 # It can be either "creator" (ZOO_CREATE_ALL_ACL) or "open" (ZOO_OPEN_ACL_UNSAFE)
117 # The default value is "open" and it can be changed to "creator" if ZK security is
118 # enabled
119
120 # high-availability.zookeeper.client.acl: open
121
122 #=====
123 # Fault tolerance and checkpointing
124 #=====
125
126 # The backend that will be used to store operator state checkpoints if
127 # checkpointing is enabled. Checkpointing is enabled when
128 # execution.checkpointing.interval > 0.
129
130 #
131 # Execution checkpointing related parameters. Please refer to CheckpointConfig and
132 # ExecutionCheckpointingOptions for more details.
133
134 #
135 # execution.checkpointing.interval: 3min
136 # execution.checkpointing.externalized-checkpoint-retention: [DELETE_ON_CANCELLATION,
137 # RETAIN_ON_CANCELLATION]
138 # execution.checkpointing.max-concurrent-checkpoints: 1
139 # execution.checkpointing.min-pause: 0
140 # execution.checkpointing.mode: [EXACTLY_ONCE, AT_LEAST_ONCE]
141 # execution.checkpointing.timeout: 10min
142 # execution.checkpointing.tolerable-failed-checkpoints: 0
143 # execution.checkpointing.unaligned: false
144
145 #
146 # Supported backends are 'jobmanager', 'filesystem', 'rocksdb', or the
147 # <class-name-of-factory>.
148
149 #
150 #开启HA，使用文件系统作为快照存储
151 state.backend: filesystem
152
153
154 # Directory for checkpoints filesystem, when using any of the default bundled
```

```
145 # state backends.
146 #
147 #默认为none, 用于指定checkpoint的data files和meta data存储的目录
148 state.checkpoints.dir: hdfs://node1:8020/flink-checkpoints
149
150 # Default target directory for savepoints, optional.
151 #
152 #默认为none, 用于指定savepoints的默认目录
153 state.savepoints.dir: hdfs://node1:8020/flink-savepoints
154
155 # Flag to enable/disable incremental checkpoints for backends that
156 # support incremental checkpoints (like the RocksDB state backend).
157 #
158 # state.backend.incremental: false
159
160 # The failover strategy, i.e., how the job computation recovers from task failures.
161 # Only restart tasks that may have been affected by the task failure, which typically
    includes
162 # downstream tasks and potentially upstream tasks if their produced data is no longer
    available for consumption.
163
164 jobmanager.execution.failover-strategy: region
165
166 #=====
167 # Rest & web frontend
168 #=====
169
170 # The port to which the REST client connects to. If rest.bind-port has
171 # not been specified, then the server will bind to this port as well.
172 #
173 #JobManager的Web界面的端口（默认：8081）
174 rest.port: 8081
175
176 # The address to which the REST client will connect to
177 #
178 #rest.address: 0.0.0.0
179
180 # Port range for the REST and web server to bind to.
181 #
182 #rest.bind-port: 8080-8090
```

```
183
184 # The address that the REST & web server binds to
185 #
186 #rest.bind-address: 0.0.0.0
187
188 # Flag to specify whether job submission is enabled from the web-based
189 # runtime monitor. Uncomment to disable.
190
191 #web.submit.enable: false
192
193 # Flag to specify whether job cancellation is enabled from the web-based
194 # runtime monitor. Uncomment to disable.
195
196 #web.cancel.enable: false
197
198 #=====
199 # Advanced
200 #=====
201
202 # Override the directories for temporary files. If not specified, the
203 # system-specific Java temporary directory (java.io.tmpdir property) is taken.
204 #
205 # For framework setups on Yarn, Flink will automatically pick up the
206 # containers' temp directories without any need for configuration.
207 #
208 # Add a delimited list for multiple directories, using the system directory
209 # delimiter (colon ':' on unix) or a comma, e.g.:
210 #     /data1/tmp:/data2/tmp:/data3/tmp
211 #
212 # Note: Each directory entry is read from and written to by a different I/O
213 # thread. You can include the same directory multiple times in order to create
214 # multiple I/O threads against that directory. This is for example relevant for
215 # high-throughput RAIDs.
216 #
217 #临时文件
218 io.tmp.dirs: /export/server/flink/tmp
219
220 # The classloading resolve order. Possible values are 'child-first' (Flink's default)
221 # and 'parent-first' (Java's default).
```



```
222 #
223 # Child first classloading allows users to use different dependency/library
224 # versions in their application than those in the classpath. Switching back
225 # to 'parent-first' may help with debugging dependency issues.
226 #
227 # classloader.resolve-order: child-first
228
229 # The amount of memory going to the network stack. These numbers usually need
230 # no tuning. Adjusting them may be necessary in case of an "Insufficient number
231 # of network buffers" error. The default min is 64MB, the default max is 1GB.
232 #
233 # taskmanager.memory.network.fraction: 0.1
234 # taskmanager.memory.network.min: 64mb
235 # taskmanager.memory.network.max: 1gb
236
237 #=====
238 # Flink Cluster Security Configuration
239 #=====
240
241 # Kerberos authentication for various components - Hadoop, ZooKeeper, and connectors -
242 # may be enabled in four steps:
243 # 1. configure the local krb5.conf file
244 # 2. provide Kerberos credentials (either a keytab or a ticket cache w/ kinit)
245 # 3. make the credentials available to various JAAS login contexts
246 # 4. configure the connector to use JAAS/SASL
247
248 # The below configure how Kerberos credentials are provided. A keytab will be used
249 # instead of
250 # a ticket cache if the keytab path and principal are set.
251
252 # security.kerberos.login.use-ticket-cache: true
253 # security.kerberos.login.keytab: /path/to/kerberos/keytab
254 # security.kerberos.login.principal: flink-user
255
256 # The configuration below defines which JAAS login contexts
257
258 # security.kerberos.login.contexts: Client,KafkaClient
259
260 #=====
261 # ZK Security Configuration
```

```
261 #=====
262
263 # Below configurations are applicable if ZK ensemble is configured for security
264
265 # Override below configuration to provide custom ZK service name if configured
266 # zookeeper.sasl.service-name: zookeeper
267
268 # The configuration below must match one of the values set in
    "security.kerberos.login.contexts"
269 # zookeeper.sasl.login-context-name: Client
270
271 #=====
272 # HistoryServer
273 #=====
274
275 # The HistoryServer is started and stopped via bin/historyserver.sh (start|stop)
276
277 # Directory to upload completed jobs to. Add this directory to the list of
278 # monitored directories of the HistoryServer as well (see below).
279 #jobmanager.archive.fs.dir: hdfs:///completed-jobs/
280
281 # The address under which the web-based HistoryServer listens.
282 #historyserver.web.address: 0.0.0.0
283
284 # The port under which the web-based HistoryServer listens.
285 #historyserver.web.port: 8082
286
287 # Comma separated list of directories to monitor for completed jobs.
288 #historyserver.archive.fs.dir: hdfs:///completed-jobs/
289
290 # Interval in milliseconds for refreshing the monitored directories.
291 #historyserver.archive.fs.refresh-interval: 10000
292
293 #blob存储文件是在群集中分发Flink作业所必需的
294 blob.storage.directory: /export/server/flink/tmp
295
296
297
```




- 配置worker

```
1 node1
2 node2
3 node3
```

- 配置mater文件

```
1 node1:8081
2 node2:8081
```

- 下载jar包到flink\lib

 flink-shaded-hadoop-2-uber-2.8.3-10.0.jar	41.3 MB	JAR 文件	2022/03/20 16:20
 flink-shaded-hadoop-3-3.1.1.7.2.9.0-173-9.0.jar	37.6 MB	JAR 文件	2022/03/20 16:19
 flink-shaded-hadoop-3-3.1.1.7.0.3.0-79-7.0.jar	33 MB	JAR 文件	2022/03/20 16:15

- 分发

```
1 scp -r /export/server/flink node2:/export/server
2 scp -r /export/server/flink node3:/export/server
3
```

- HA高可用，设置HDFS上的路径用于保存ha的数据，防止出现当前集群jobmanager挂掉恢复最新状态
- 需要先开启zookeeper，再启动flink集群，通过start-cluster.sh
- 切换jobmanager实现HA高可用
  - 关闭node1上的 jobmanager 进程
  - 查看 node2 上 8081 web的log日志，查看是否 granted leadership

## YARN部署

- 配置
  - yarn-site.xml 中修改一下memcheck 置为 false ，不让检查内存是否可用。

```
1 <property>
2     <name>yarn.nodemanager.vmem-check-enabled</name>
3     <value>>false</value>
4 </property>
```

- 然后分发到各个节点上
- yarn-session + flink run
  - 应用场景：大量的小任务，当小任务执行完毕之后并不会关闭session，小任务之间共享session（内存和CPU cores）不隔离资源。

```
1 # 开启 yarn-session 会话
```

```
2 yarn-session.sh -tm 1024 -s 2 -d
3 # -tm taskmanager 的内存大小
4 # -s slot 数
5 # -d daemon 后台执行
6 flink run -p 2 /export/server/flink/examples/batch/WordCount.jar
```

- kill掉一直运行的session

```
1 yarn application -kill application_1638083192874_0001
```

- 每个任务都是直接 flink run 执行 per-job
  - 应用场景： 适合于大多数生产环境的，任务的执行，每个任务一个session，程序执行完毕关闭会话。

```
1 flink run \  
2 -m yarn-cluster \  
3 -yjm 1024m \  
4 -ytm 1024m \  
5 /export/server/flink/examples/batch/WordCount.jar \  
6 --input hdfs://node1:8020/words.txt
```